



# Google Cloud

---

Serverless

Messaging with  
Cloud Pub/Sub

# Agenda

---

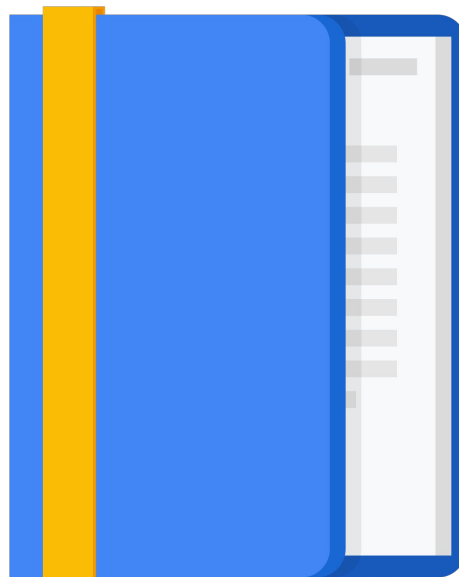
Processing Streaming Data

Cloud Pub/Sub

Cloud Dataflow Streaming  
Features

BigQuery and Bigtable Streaming  
Features

Advanced BigQuery Functionality



# Cloud Pub/Sub



Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:

Availability  
Durability  
Scalability

100s of  
milliseconds

# Cloud Pub/Sub



Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:

Availability  
Durability  
Scalability

100s of  
milliseconds

# Cloud Pub/Sub



Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:

Availability  
Durability  
Scalability

100s of  
milliseconds

# Cloud Pub/Sub




Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:


Availability  
Durability  
Scalability

100s of  
milliseconds

# Example of a Cloud Pub/Sub application

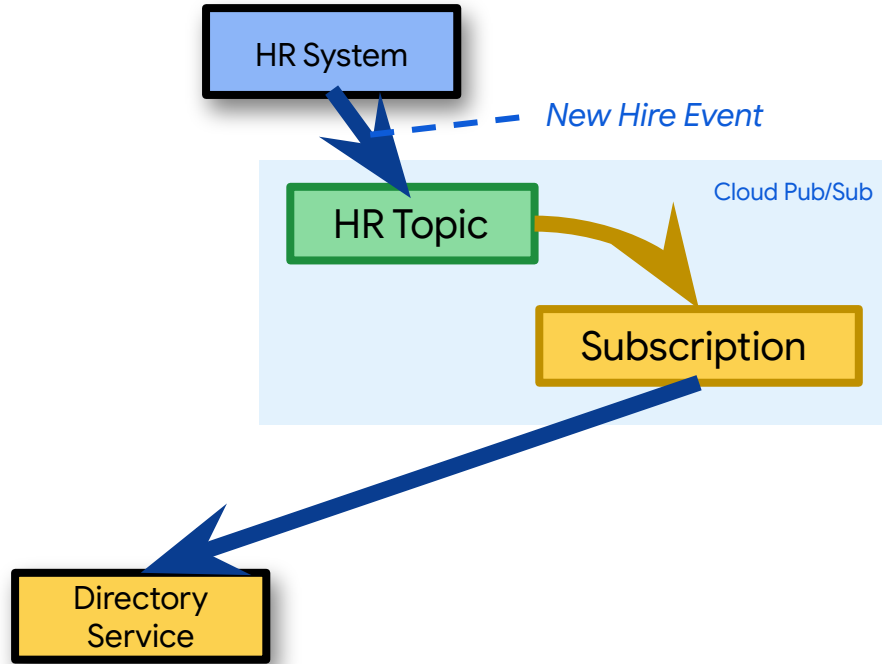


A new employees arrives causing a new hire event.

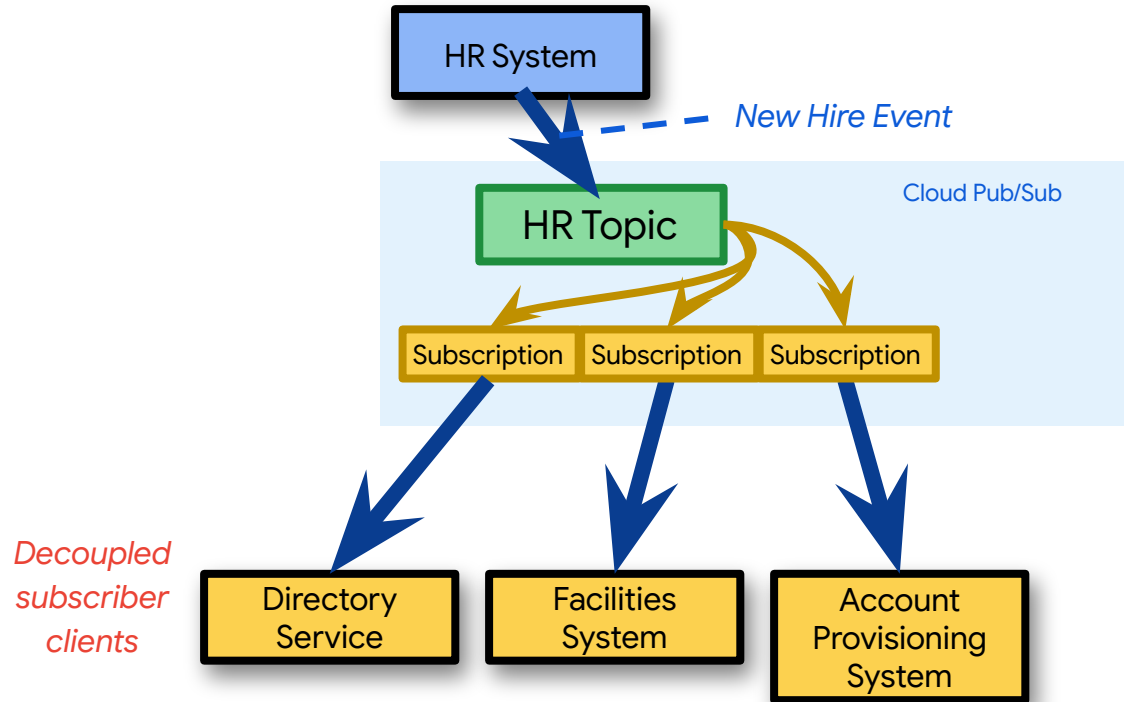




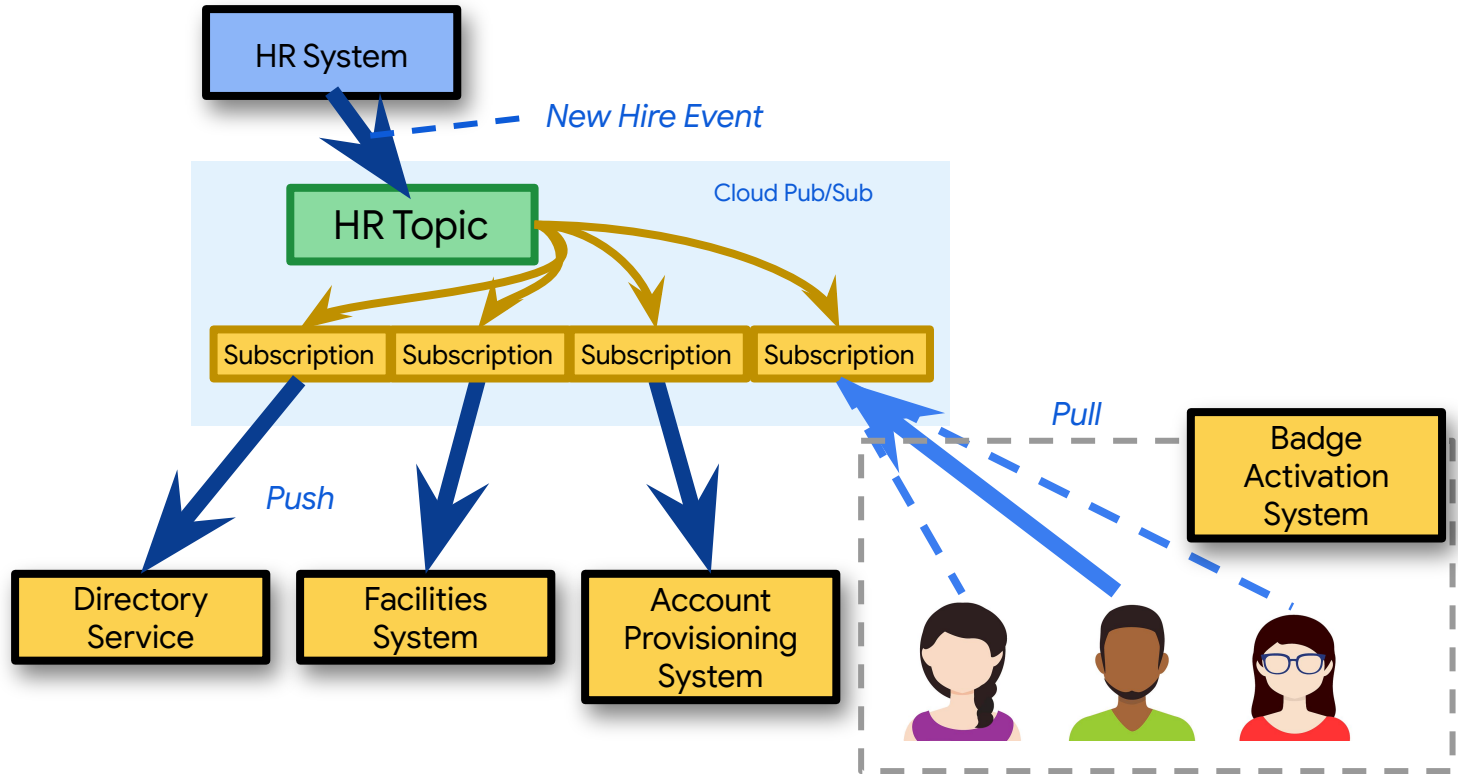
# The message is sent from Topic to Subscription



# There can be multiple Subscriptions for each Topic

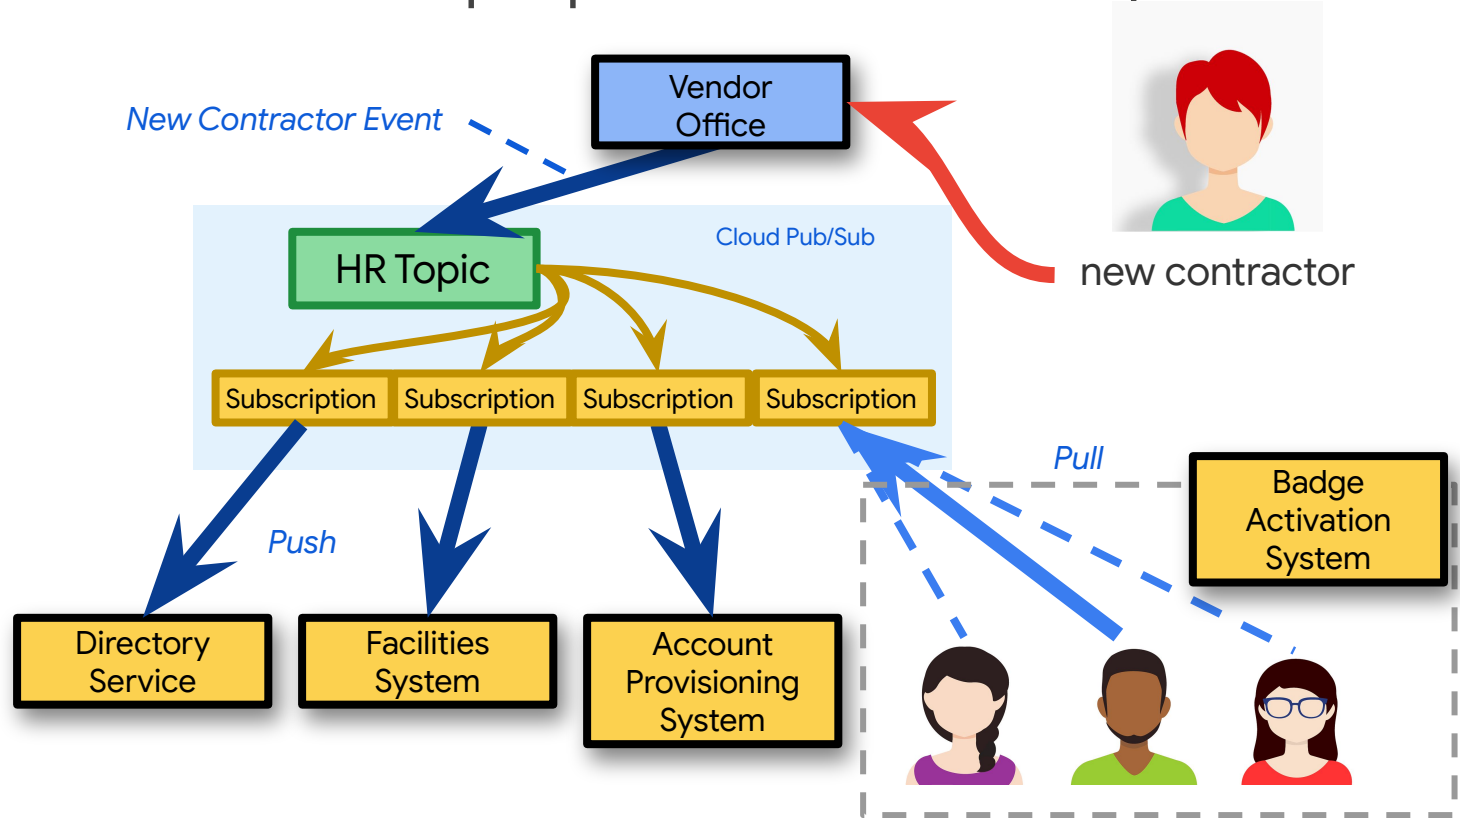


# And there can be multiple subscribers per Subscription

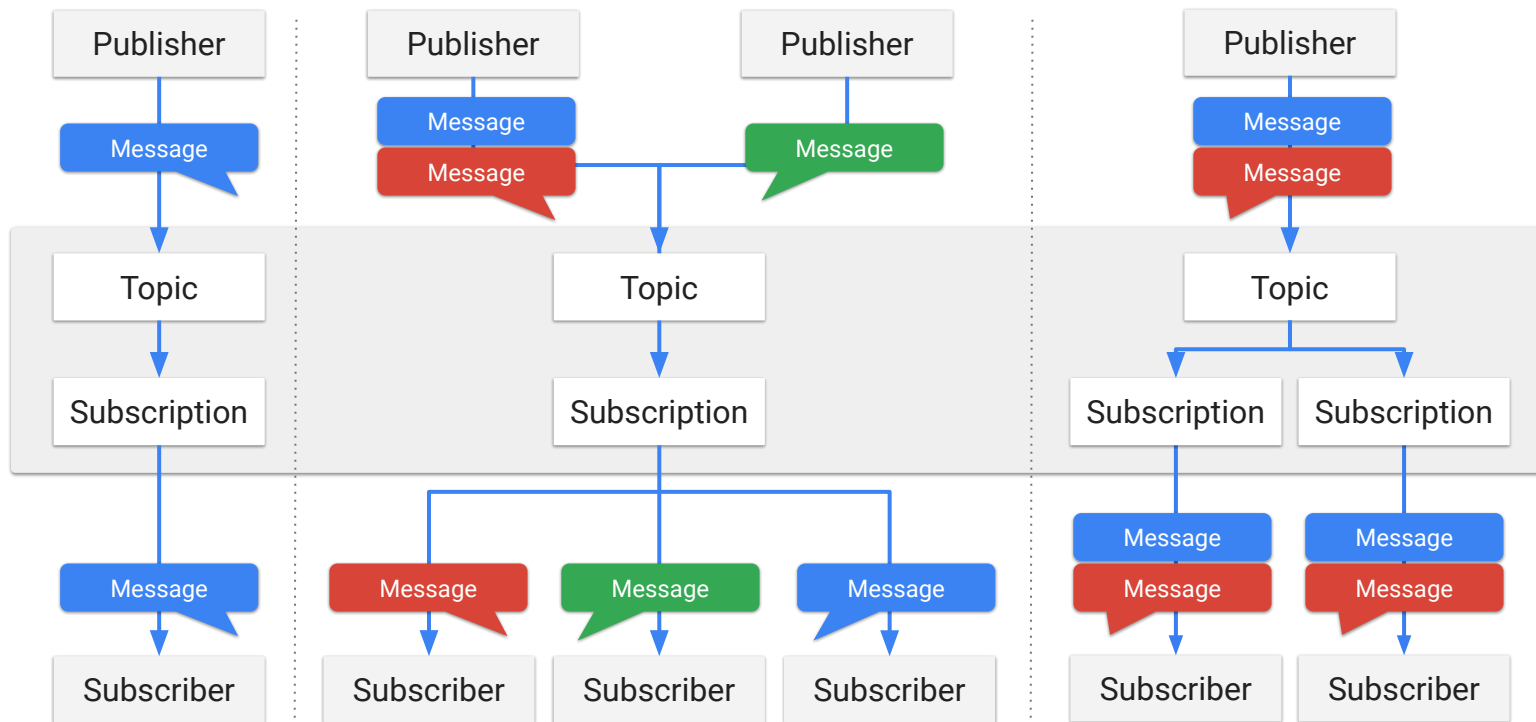


# And there can be multiple publishers to the Topic

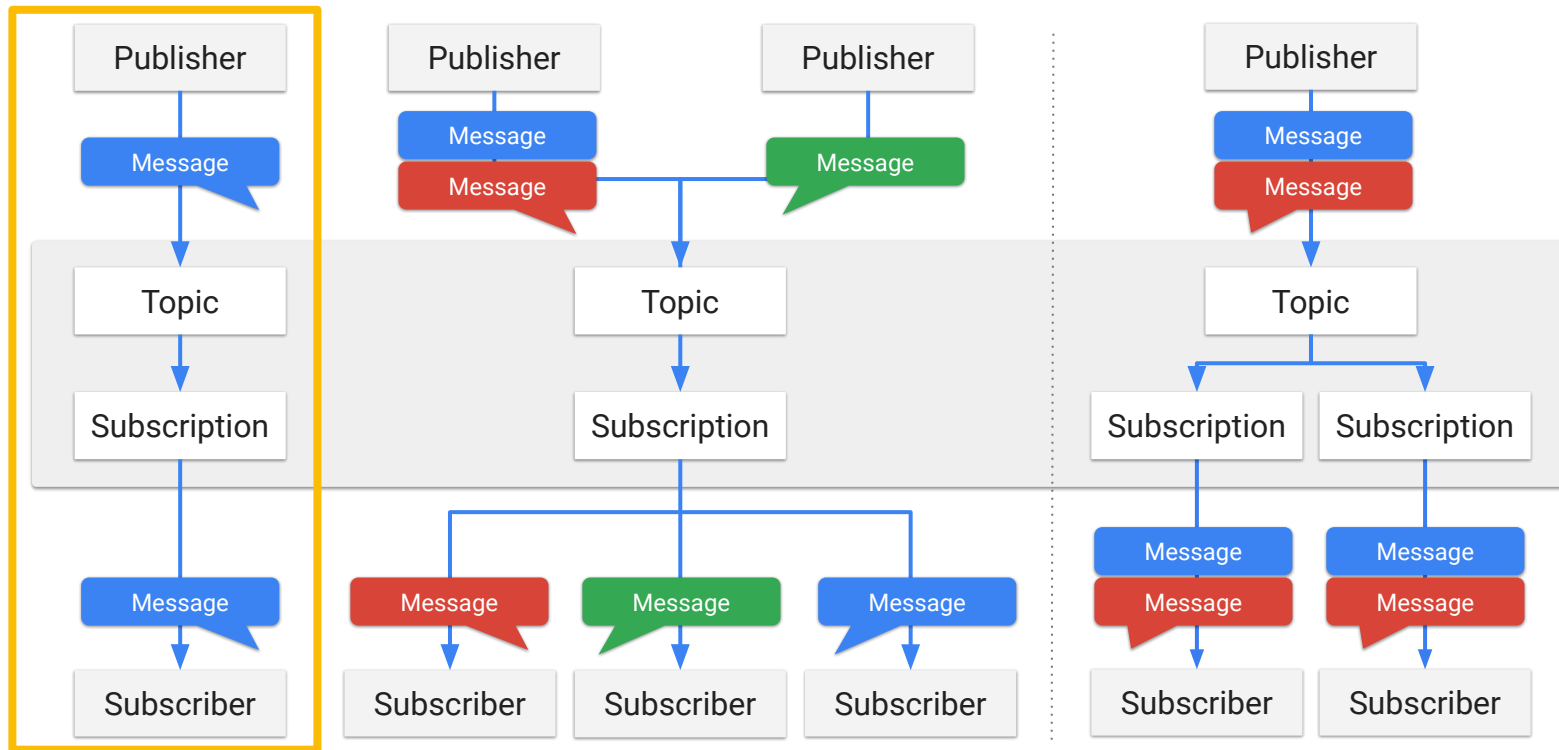
*Decoupled  
publisher  
clients*



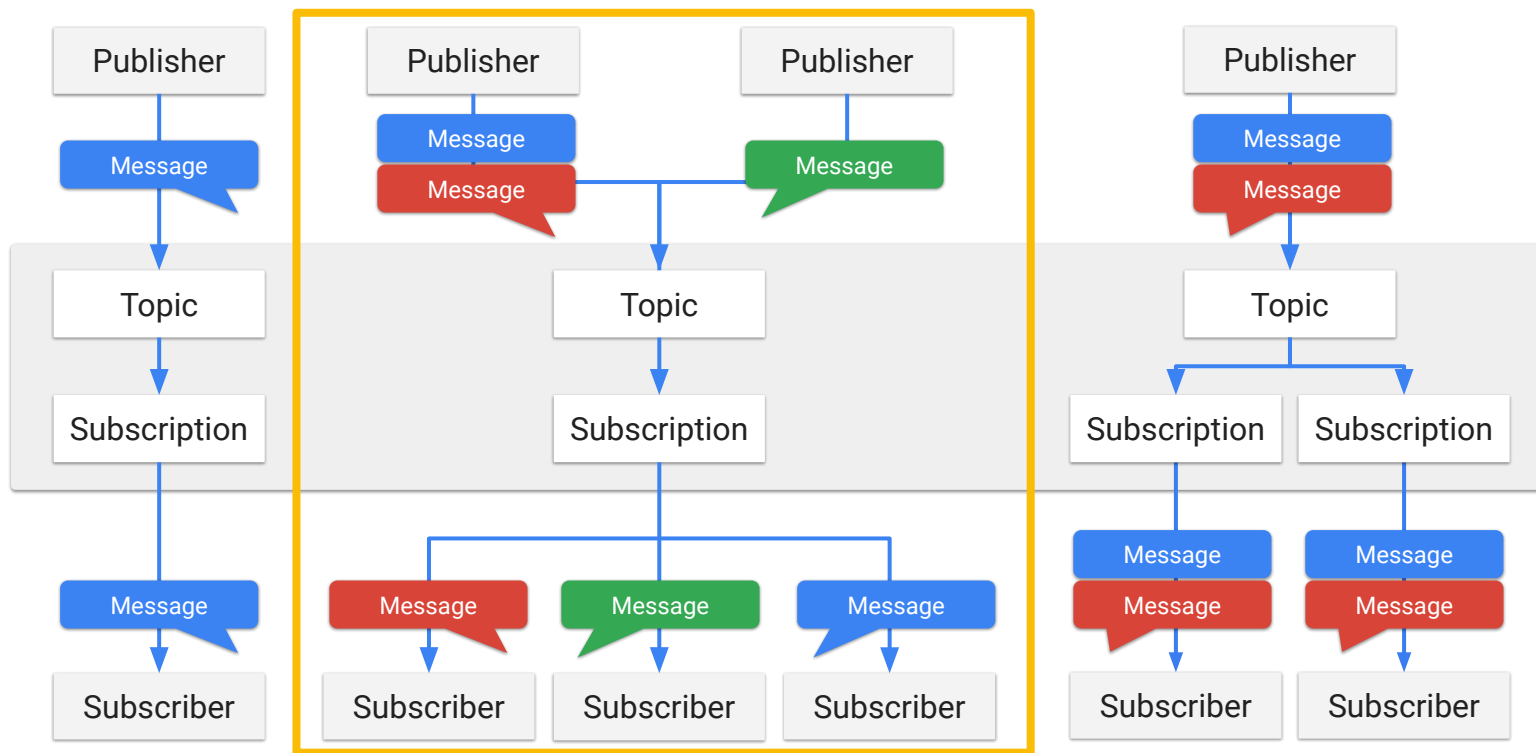
# Publish/Subscribe patterns



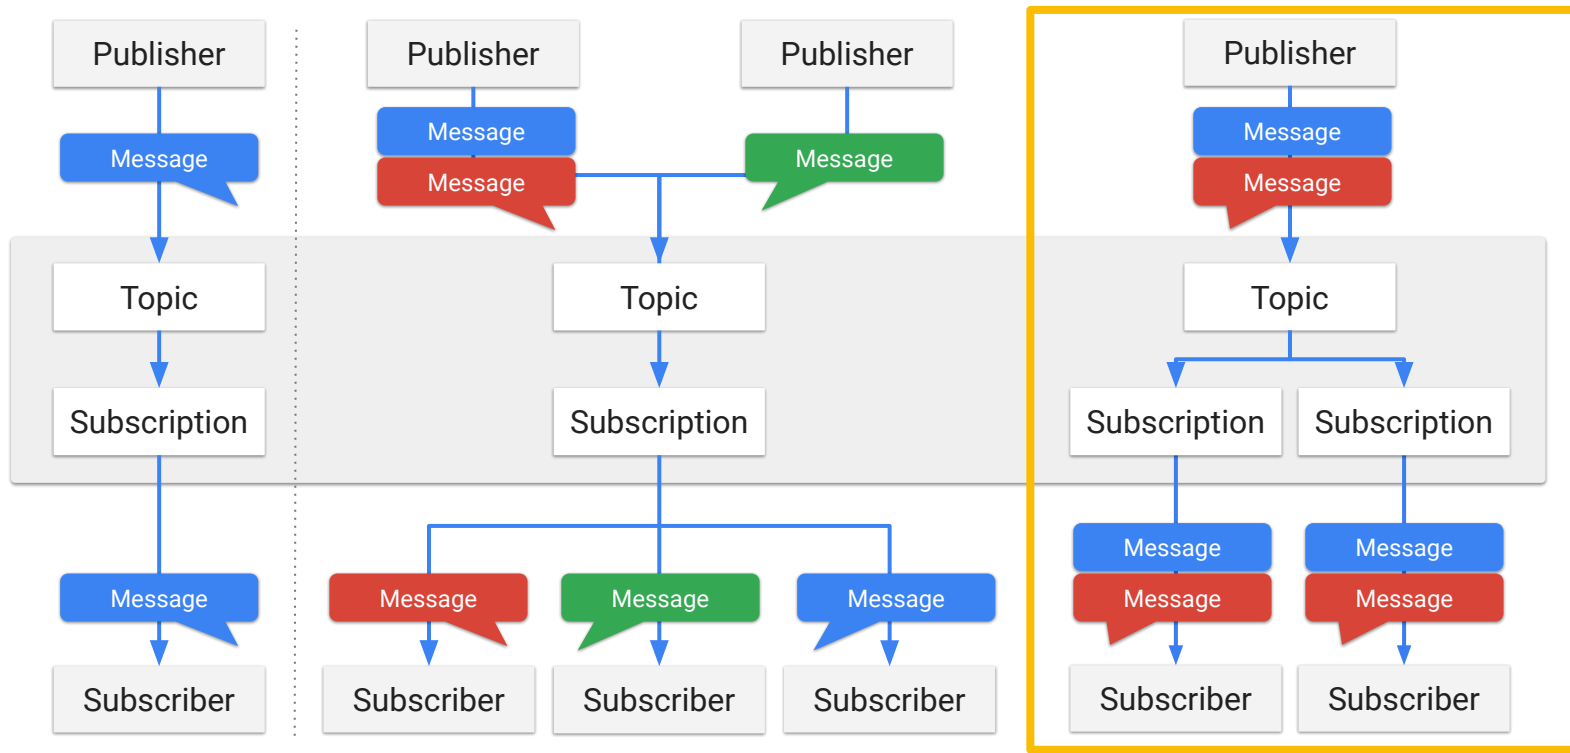
# Publish/Subscribe patterns



# Publish/Subscribe patterns





# Publish/Subscribe patterns



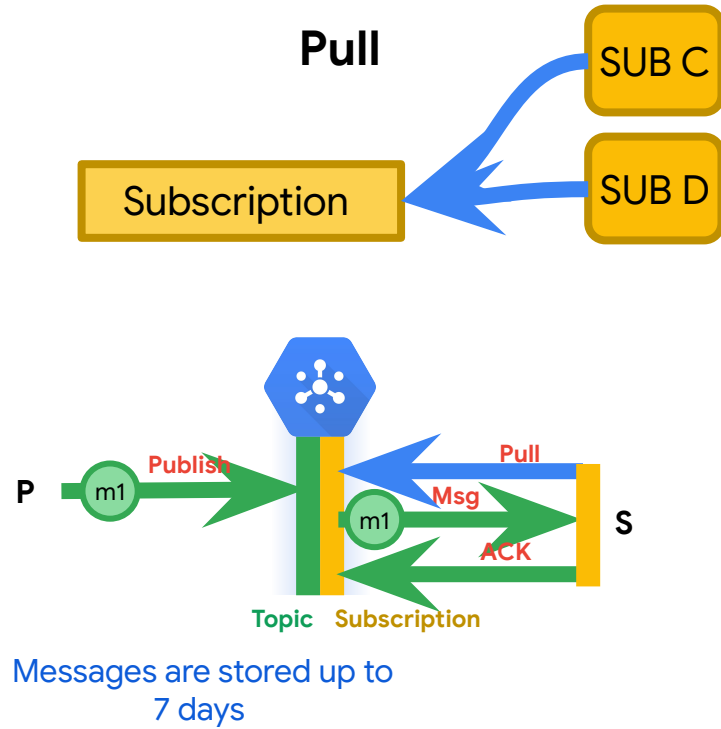
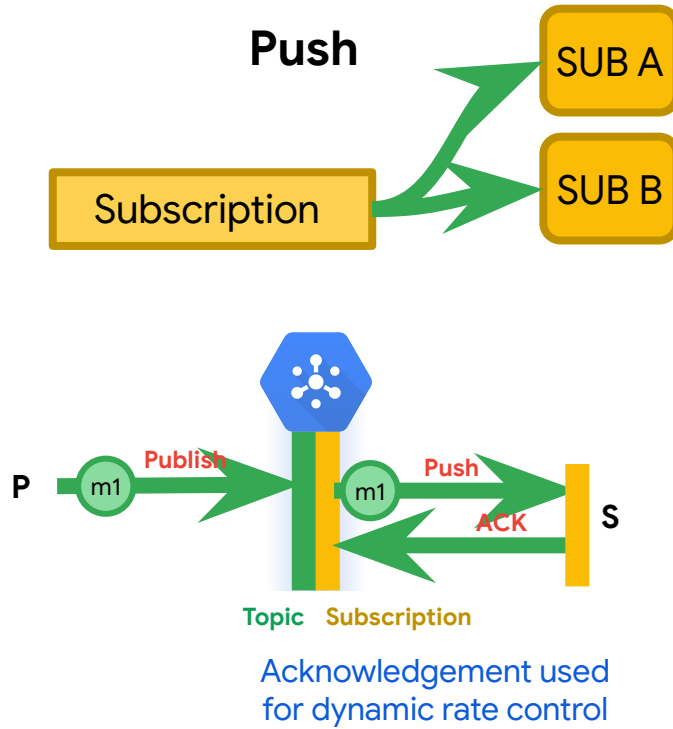


# Cloud Pub/Sub provides both Push and Pull delivery



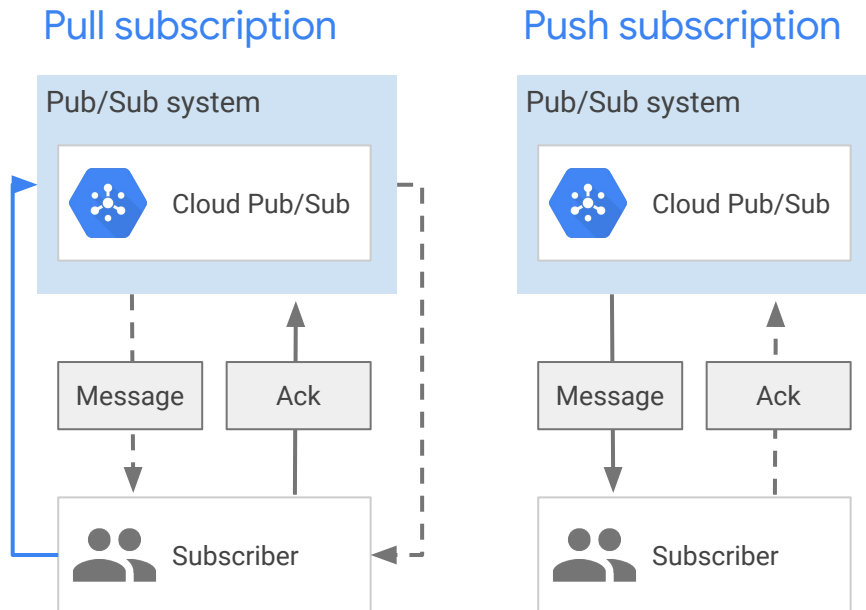
Messages are stored up to  
7 days

# Cloud Pub/Sub provides both Push and Pull delivery

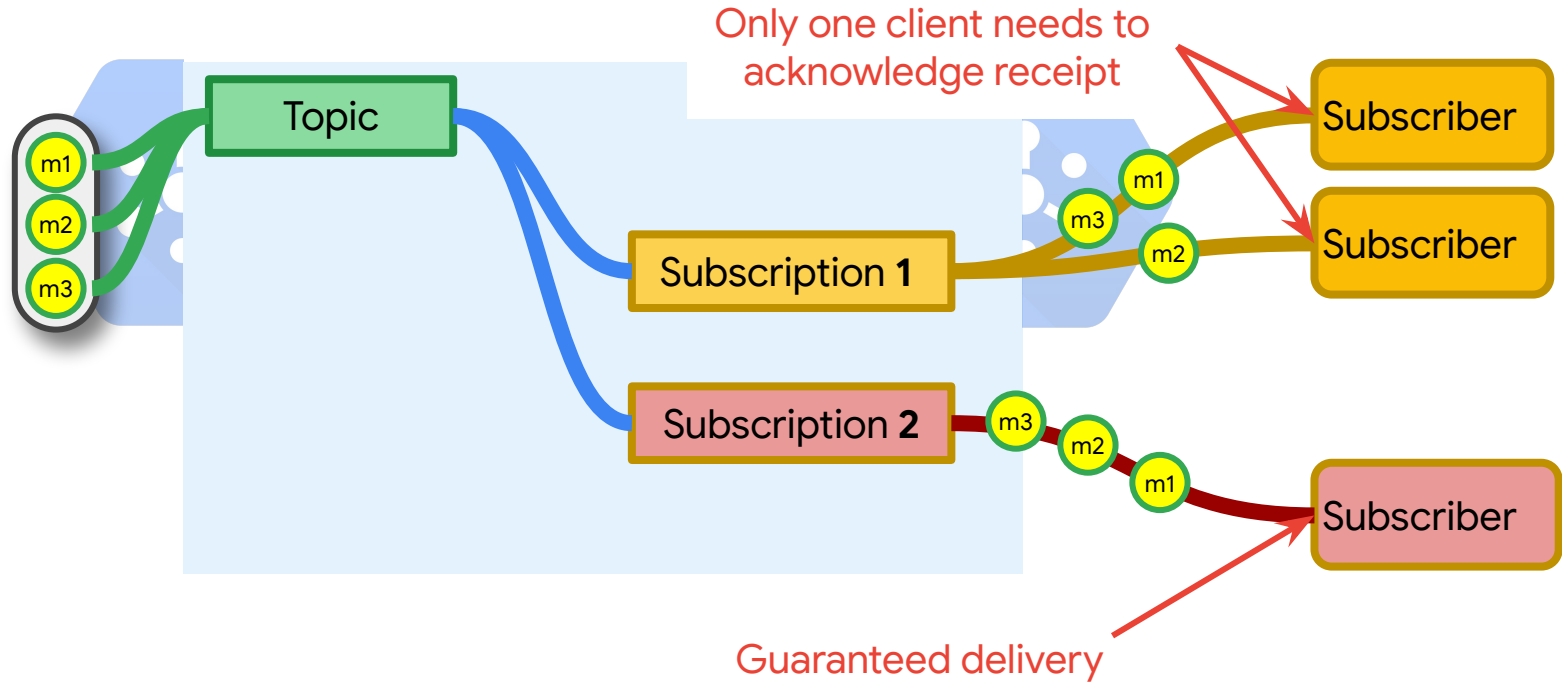


# At least once delivery guarantee

- A subscriber ACKs each message for every subscription
- A message is resent if subscriber takes more than **ackDeadline** to respond
- Messages are stored for up to 7 days
- A subscriber can extend the deadline per message



# Subscribers can work as a team or separately



# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

[Create topic](#)

# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

← -- Set topic name

↑  
Message

↑  
Send attribute

# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

← -- Set topic name

↑  
Message

↑  
Send attribute



# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Set topic name

Message

Send attribute

Create a client

# Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method  
Callback function

# Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method  
Callback function

# Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method  
Callback function

# Subscribing with Cloud Pub/Sub using synchronous pull

```
gcloud pubsub subscriptions create --topic sandiego mySub1
```

Create subscription

```
gcloud pubsub subscriptions pull --auto-ack mySub1
```

Pull subscription

```
import time

from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()

subscription_path = subscriber.subscription_path(project_id, subscription_name)

NUM_MESSAGES = 2
ACK_DEADLINE = 30
SLEEP_TIME = 10

# The subscriber pulls a specific number of messages.
response = subscriber.pull(subscription_path, max_messages=NUM_MESSAGES)
```

Set subscription name

Create a client

``projects/{project_id}/subscriptions/{subscription_name}`` ← subscription\_path format

Subscriber is non-blocking

Keep the main thread from exiting to allow it to process messages synchronously

By default, the Publisher batches messages; turn this off if you desire lower latency



Batching messages: throughput versus latency

# Changing the batch settings in Cloud Pub/Sub

Python

```
from google.cloud import pubsub
from google.cloud.pubsub import types

client = pubsub.PublisherClient(
    batch_settings=BatchSettings(max_messages=500),
)
```

[Change batch setting](#)

# Pub/Sub: latency, out-of-order, duplication will happen

- Latency -- no guarantees
- Messages can be delivered in any order, especially with large backlog
- Duplication may happen



# Cloud Pub/sub with Dataflow: Exactly once, ordered processing




Cloud Pub/Sub delivers at least once




Cloud Dataflow: Deduplicate, order, and window





Separation of concerns → scale



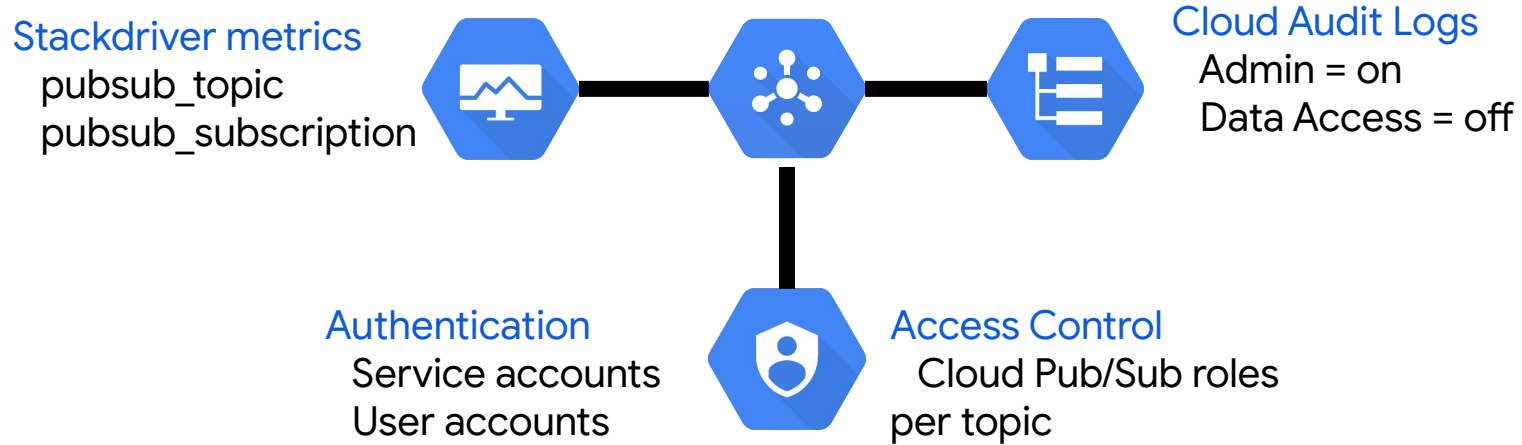
# Use Cloud Pub/Sub for streaming resilience



# Use Cloud Pub/Sub for streaming resilience



# Security, monitoring, and logging for Cloud Pub/Sub





---

Publish Streaming Data  
into Pub/Sub

# Lab Objectives

---

Create a Pub/Sub topic and subscription

Simulate your traffic sensor data into Pub/Sub