



## Cloud Dataflow Streaming Features

# Agenda

---

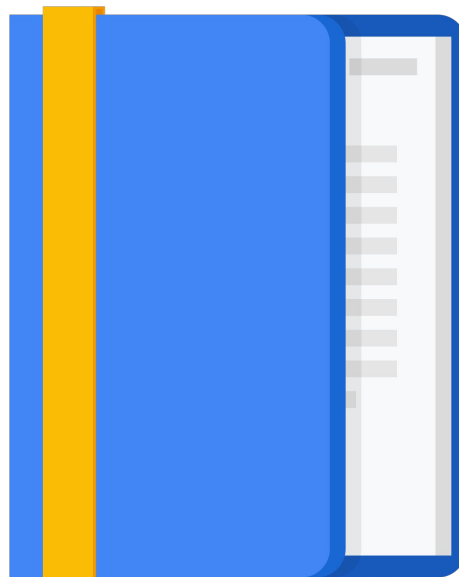
Processing Streaming Data

Cloud Pub/Sub

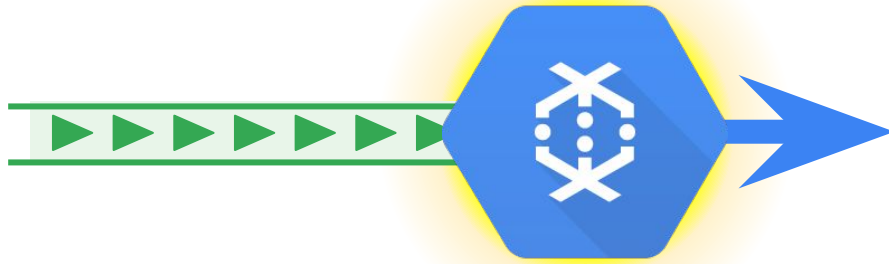
Cloud Dataflow Streaming  
Features

BigQuery and Bigtable Streaming  
Features

Advanced BigQuery Functionality



# Streaming features of Cloud Dataflow



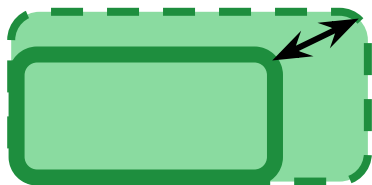
Cloud  
Dataflow

Qualities that Cloud Dataflow  
contributes to Data Engineering  
solutions:

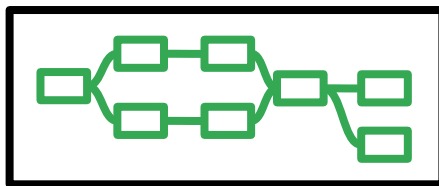
Scalability  
Low latency

# Continuing from the Data Processing course

Unbounded PCollection



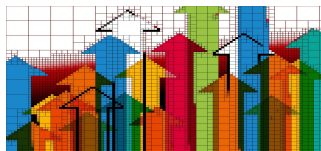
Pipeline



Streaming Jobs



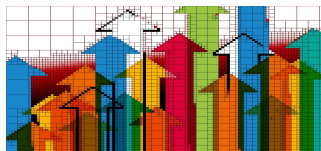
# There are challenges with processing streaming data



## **Scalability**

Streaming data generally only grows larger and more frequent

# There are challenges with processing streaming data



## **Scalability**

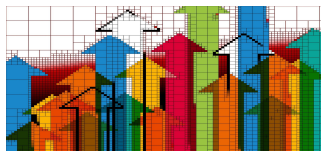
Streaming data generally only grows larger and more frequent



## **Fault Tolerance**

Maintain fault tolerance despite increasing volumes of data

# There are challenges with processing streaming data



## **Scalability**

Streaming data generally only grows larger and more frequent



## **Fault Tolerance**


Maintain fault tolerance despite increasing volumes of data



## **Model**

Is it streaming or repeated batch?

# There are challenges with processing streaming data



## **Scalability**

Streaming data generally only grows larger and more frequent



## **Fault Tolerance**

Maintain fault tolerance despite increasing volumes of data



## **Model**

Is it streaming or repeated batch?



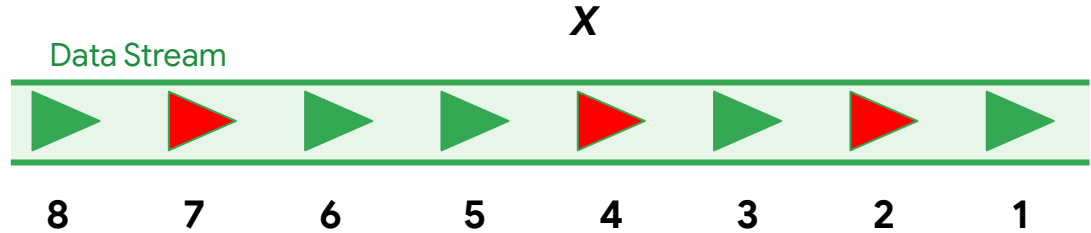
## **Timing**

What if data arrives late?

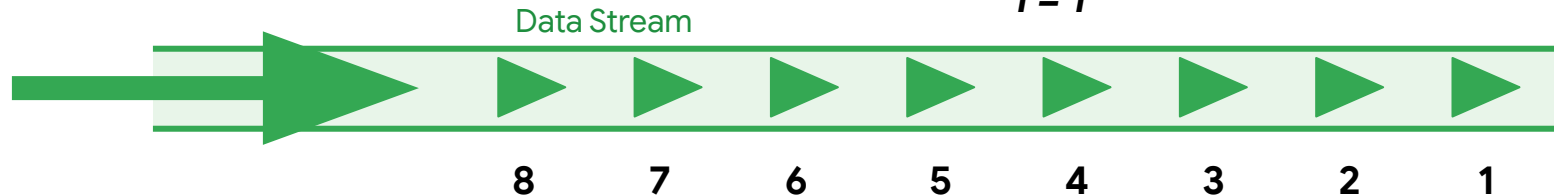


# How do you aggregate an unbounded set?

Filtering is straightforward



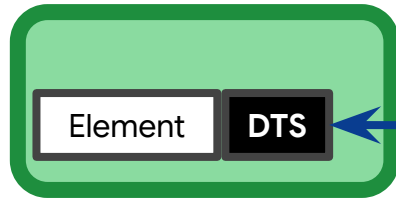
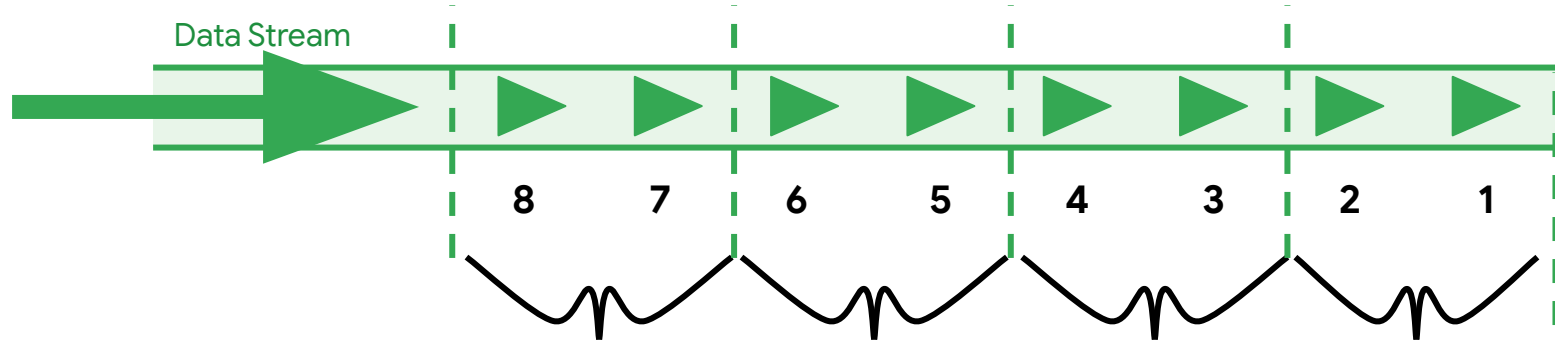
But how do you take an average on an unbounded stream of data?



$$\sum_{i=1}^n x_i$$

What is the stopping limit?

# Divide the stream into a series of finite windows




Windows are defined by element date-timestamp (metadata, usually Pub/Sub ingest time)

$$\sum_{i=w}^n x_i$$

$$\sum_{i=w}^n x_i$$

$$\sum_{i=w}^n x_i$$

# Modify the date-timestamp with a PTransform if needed



## Code to modify date-timestamp

Python

```
yield beam.window.TimestampedValue(element, unix_timestamp)
```

Java

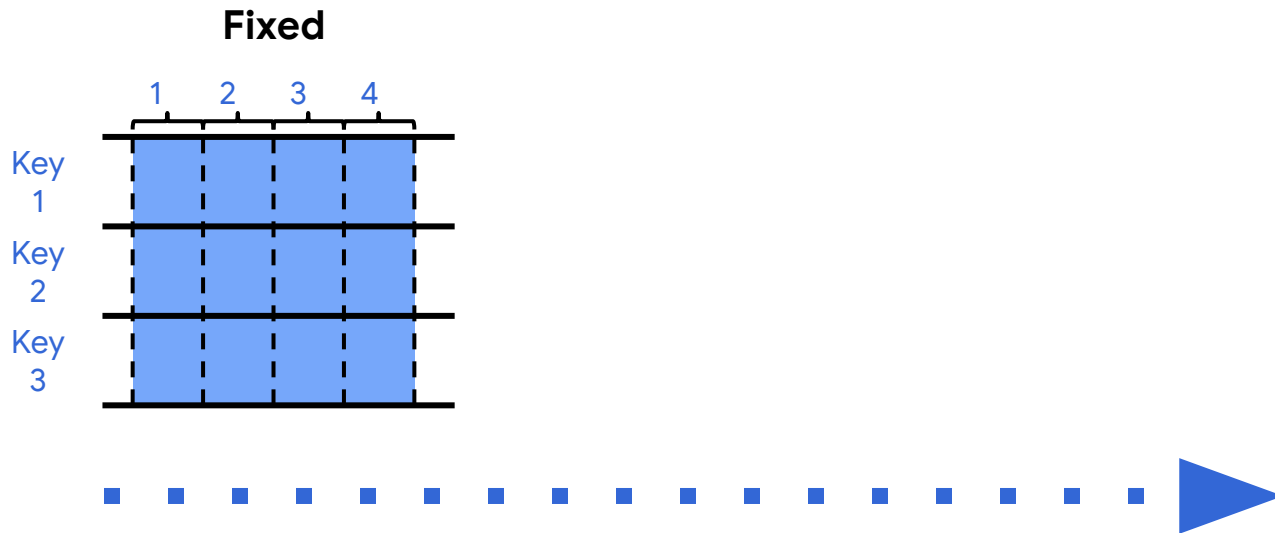
```
c.outputWithTimestamp (element, timestamp);
```

# Cloud Dataflow Windowing

# Three kinds of windows fit most circumstances


- Fixed
- Sliding
- Sessions

# Three kinds of windows fit most circumstances



Windowing divides data into  
time-based finite chunks

# Three kinds of windows fit most circumstances




Windowing divides data into  
time-based finite chunks

Often required when doing  
aggregations over unbounded data



# Three kinds of windows fit most circumstances



Windowing divides data into time-based finite chunks

Often required when doing aggregations over unbounded data

# Setting time windows

## Fixed-time windows

```
from apache_beam import window
fixed_windowed_items = (
    items | 'window' >> beam.WindowInto(window.FixedWindows(60)))
```

Python

## Sliding time windows

```
from apache_beam import window
sliding_windowed_items = (
    items | 'window' >> beam.WindowInto(window.SlidingWindows(30, 5)))
```

Python

## Session windows

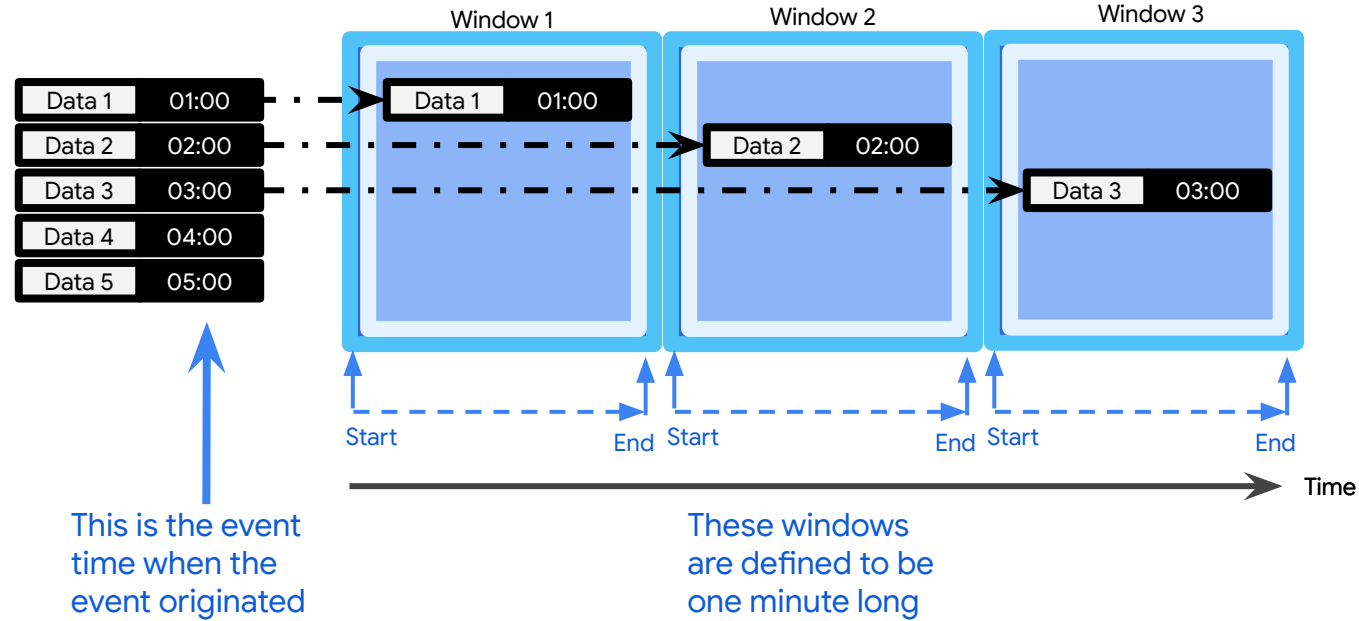
```
from apache_beam import window
session_windowed_items = (
    items | 'window' >> beam.WindowInto(window.Sessions(10 * 60)))
```

Python

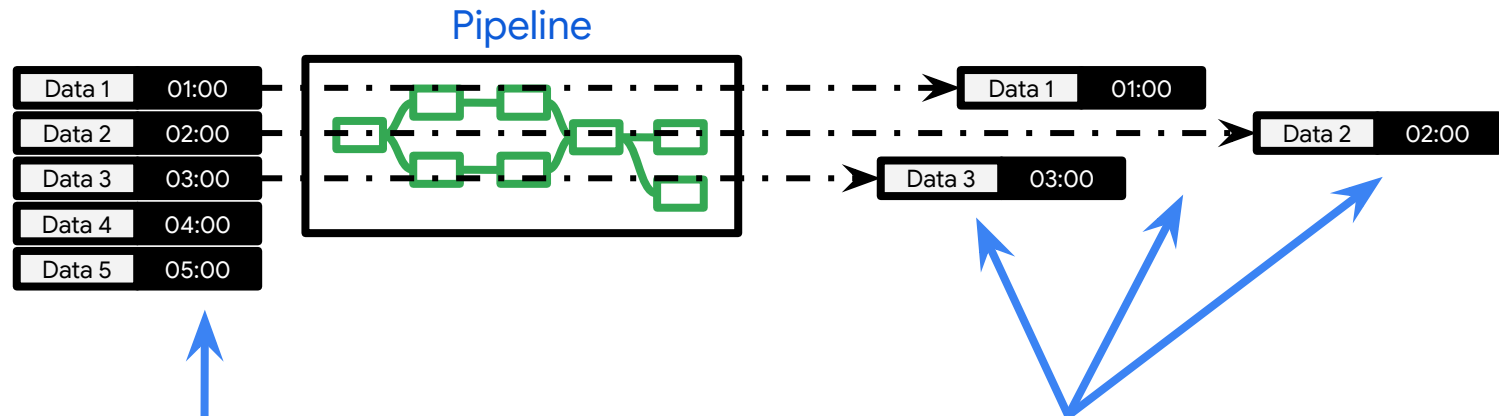
### Remember:

you can apply windows to batch data, although you may need to generate the metadata date-timestamp on which windows operate.

# Windowing by time if there is no latency



# Pipeline processing can introduce latency



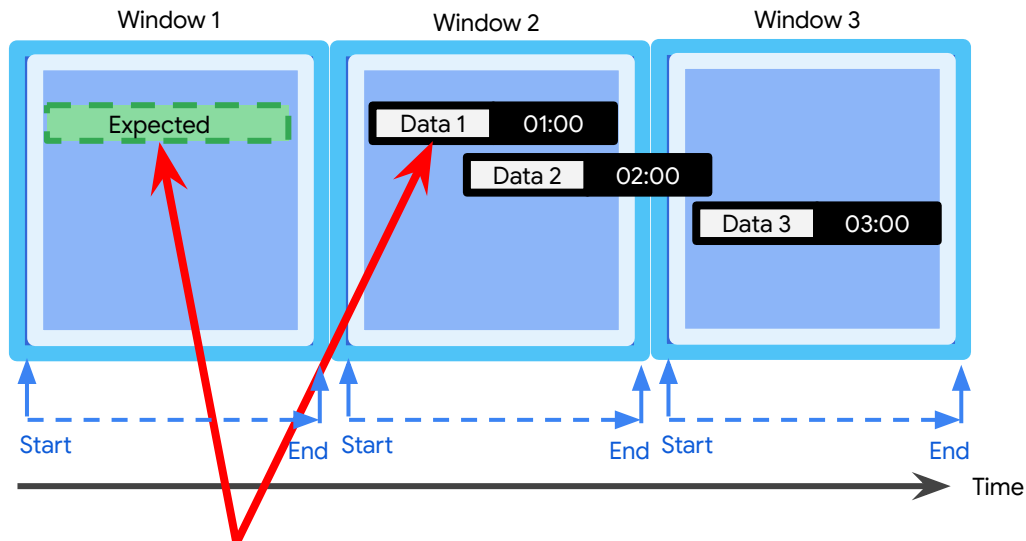
This is the event time when the event originated

The time relationships at origination are not preserved.

They are arriving later than anticipated. And some of them are outside the original 1 minute window.

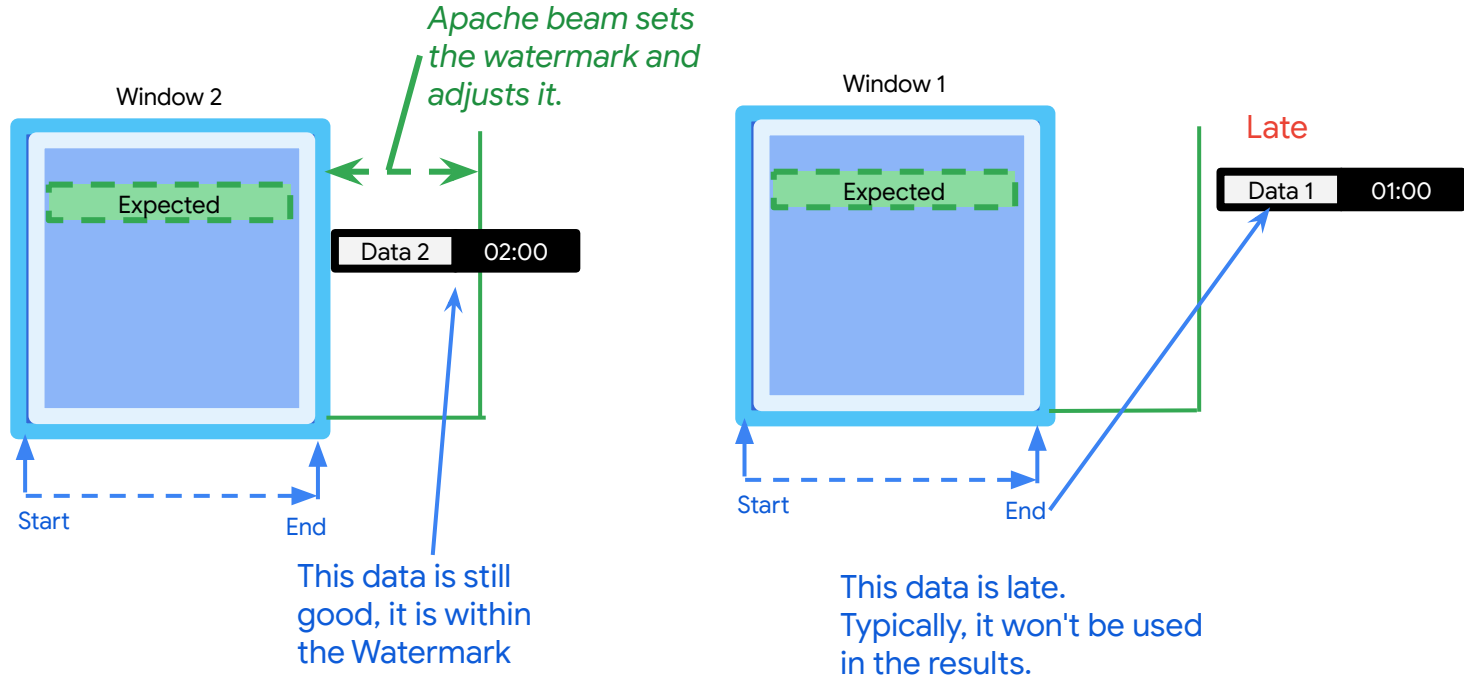
# How should Cloud Dataflow deal with this situation?

The data could be a little past the window or a lot. Data 2 is a little outside of Window 2. Data 1 is completely outside of Window 1.

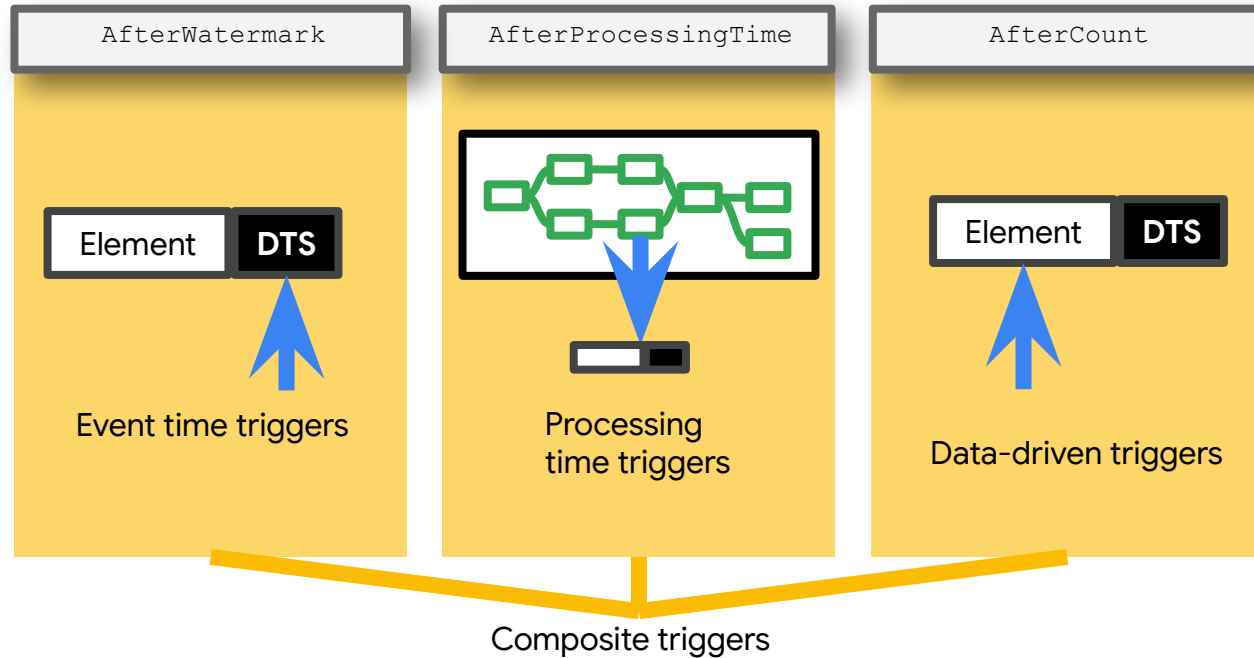


The difference in time from when data was expected to when it actually arrived is called the **lag time**.

# Watermarks provide flexibility for a little lag time



The default is to trigger at the watermark, but we can also add custom trigger(s)



# Some example triggers

```
pcollection | WindowInto(  
    SlidingWindows(60, 5),  
    trigger=AfterWatermark(  
        early=AfterProcessingTime(delay=30),  
        late=AfterCount(1))  
    accumulation_mode=AccumulationMode.ACCUMULATING)
```

# Sliding window of 60 seconds, every 5 seconds  
# Relative to the watermark, trigger:  
# -- fires 30 seconds after pipeline commences  
# -- and for every late record (< allowedLateness)  
# the pane should have all the records

```
pcollection | WindowInto(  
    FixedWindows(60),  
    trigger=Repeatedly(  
        AfterAny(  
            AfterCount(100),  
            AfterProcessingTime(1 * 60))),  
    accumulation_mode=AccumulationMode.DISCARDING)
```

# Fixed window of 60 seconds  
# Set up a composite trigger that triggers ...  
# whenever either of these happens:  
# -- 100 elements accumulate  
# -- every 60 seconds (ignore watermark)  
# the trigger should be with only new records



# You can allow late data past the watermark in Java

## Allowing Late Data

```
PCollection<String> items = ...;
```

Java

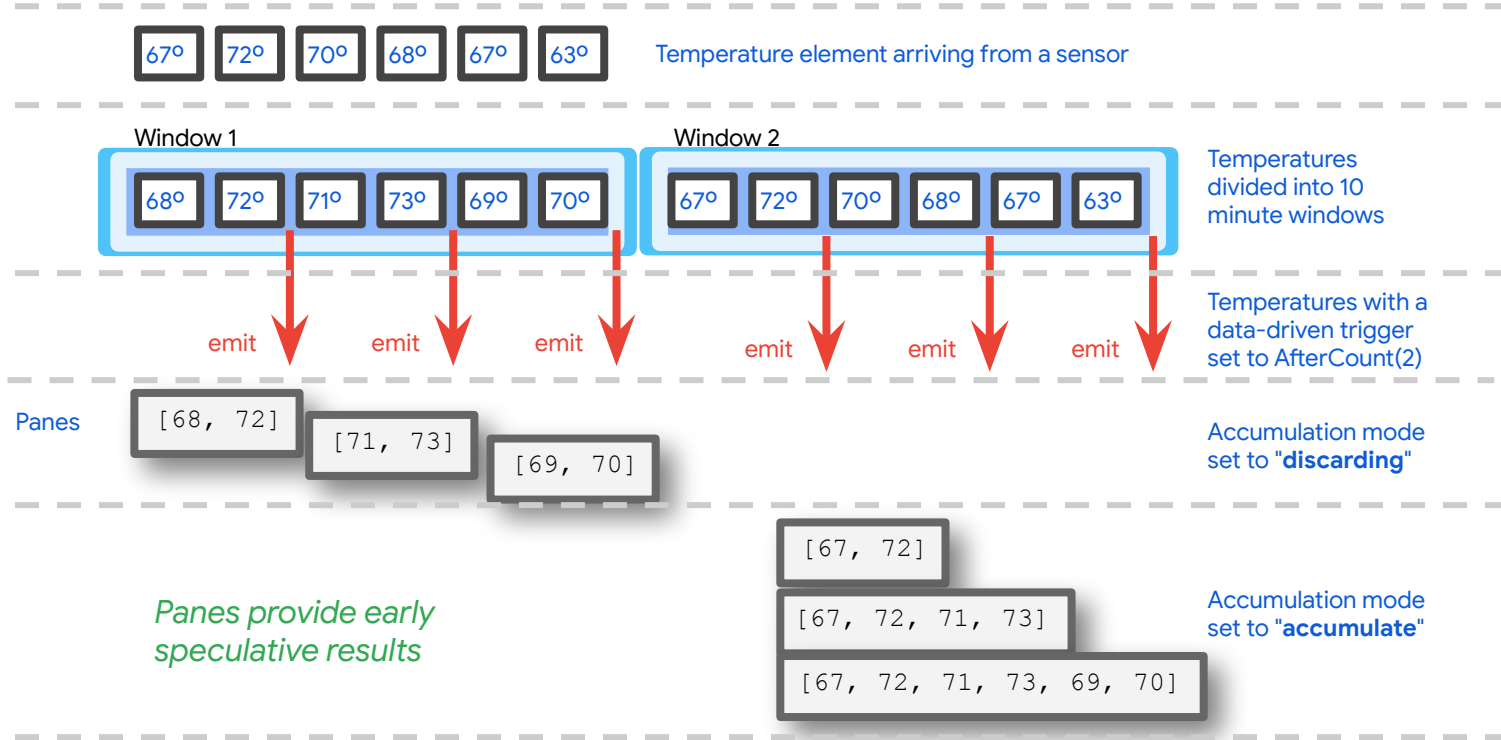
```
PCollection<String> fixedWindowedItems = items.apply(  
    Window.<String>into(FixedWindows.of(Duration.standardMinutes(1)))  
    .withAllowedLateness(Duration.standardDays(2)));
```

Python

**Not supported for Python at this time.**

That means 100% of late data is just discarded. And features designed to do something with late data simply do not work in a Python pipeline.

# Accumulation modes: what to do with additional events





---

# Streaming Data Pipelines

## Objectives

- Launch Dataflow and run a Dataflow job
- Understand how data elements flow through the transformations of a Dataflow pipeline
- Connect Dataflow to Pub/Sub and BigQuery
- Observe and understand how Dataflow autoscaling adjusts compute resources to process input data optimally
- Learn where to find logging information created by Dataflow
- Explore metrics and create alerts and dashboards with Stackdriver Monitoring